

# **Assignment 3: Classification & Model Selection**

***CS 589 - ML***

Shubham Shetty (shubhamshett@umass.edu)  
Brinda Murulidhara (bmurulidhara@umass.edu)  
Adarsh Kolya (akolya@umass.edu)

*March 2021*

# Preface

Before running our codes, the following packages are imported and test & training data are assigned to variables. Also, some reusable functions are defined -

```
1 # Import Statements
2 import numpy as np
3 import sklearn as sk
4 from sklearn.tree import *
5 from sklearn.linear_model import *
6 from sklearn.svm import *
7 from sklearn.metrics import *
8 from sklearn.neighbors import *
9 from sklearn.model_selection import KFold
10 from datetime import datetime
11 import math
12 import csv
13 import pandas as pd
14 import sys
15 import matplotlib.pyplot as plt
16 import multiprocessing
17
18 # Load Data
19 data = np.load("data.npz")
20
21 X_trn = data["X_trn"]
22 y_trn = data["y_trn"]
23 X_tst = data["X_tst"]
24
25
26 # Common Functions
27
28 # Show image from data
29 def show(x):
30     img = x.reshape((3,32,32)).transpose(1,2,0)
31     plt.imshow(img)
32     plt.axis('off')
33     plt.draw()
34     plt.pause(0.01)
35
36 # Write data to CSV file on local system
37 def write_csv(y_pred, filename):
38     """Write a 1d numpy array to a Kaggle-compatible .csv file"""
39     with open(filename, 'w+') as csv_file:
40         csv_writer = csv.writer(csv_file)
41         csv_writer.writerow(['Id', 'Category'])
42         for idx, y in enumerate(y_pred):
43             csv_writer.writerow([idx, y])
44
45 # Helper function to print table
46 # Input variable table is a numpy array and headers is an array of
47 # column headers
47 def prettyPrintTable(table, headers) :
```

```
48     for i, d in enumerate(table):
49         if i == 0 :
50             line = '|'.join(str(x).ljust(30) for x in headers)
51             print(line)
52             print('-' * len(line))
53
54         line = '|'.join(str(x).ljust(30) for x in d)
55         print(line)
56     print()
```

# Answer 1

Given -

$$V = \{0, 1\}$$

Probabilities for training outputs are-

$$P(0) = 1/4$$

$$P(1) = 3/4$$

Let  $D$  be cross entropy and  $I$  represent information gain. Then, cross entropy before the split is (assuming log base  $e$ )-

$$D_0 = -[\frac{1}{4}\log(\frac{1}{4}) + \frac{3}{4}\log(\frac{3}{4})] = 0.562$$

1. Split at  $x = 0.5$

$$D_{m < 0.5} = 0$$

$$D_{m \geq 0.5} = -[\frac{1}{4}\log(\frac{1}{4}) + \frac{3}{4}\log(\frac{3}{4})] = 0.562$$

Therefore,

$$I = 0.562 - 0.562 = 0$$

2. Split at  $x = 1.5$

$$D_{m < 1.5} = -\log(1) = 0$$

$$D_{m \geq 1.5} = -[\frac{2}{3}\log(\frac{2}{3}) + \frac{1}{3}\log(\frac{1}{3})] = 0.636$$

Therefore,

$$I = 0.562 - \frac{3}{4} * 0.634 = 0.0846$$

3. Split at  $x = 2.5$

$$D_{m < 2.5} = -[\frac{1}{2}\log(\frac{1}{2}) + \frac{1}{2}\log(\frac{1}{2})] = 0.693$$

$$D_{m \geq 2.5} = -\log(1) = 0$$

Therefore,

$$I = 0.562 - \frac{1}{2} * 0.693 = 0.2155$$

4. Split at  $x = 3.5$

$$D_{m < 3.5} = -[\frac{2}{3}\log(\frac{2}{3}) + \frac{1}{3}\log(\frac{1}{3})] = 0.636$$

$$D_{m \geq 3.5} = -\log(1) = 0$$

Therefore,

$$I = 0.562 - \frac{3}{4} * 0.634 = 0.0846$$

5. Split at  $x = 4.5$

$$D_{m < 4.5} = -[\frac{1}{4}\log(\frac{1}{4}) + \frac{3}{4}\log(\frac{3}{4})] = 0.562$$

$$D_{m \geq 4.5} = 0$$

Therefore,

$$I = 0.562 - 0.562 = 0$$

## Answer 2

Time complexity to evaluate that classification tree on a single new input-

$$\mathcal{O}(M)$$

In the worst case, the longest path from root to leaf of the tree may be traversed to evaluate the classification tree on a single new input. The longest root to leaf path is equal to the depth (M) of the tree.

## Answer 3

Time complexity to train a classification stump-

$$\mathcal{O}(DN \log N)$$

There are  $D$  dimensions, and it takes  $N \log N$  time for each sorting operation. Checking all the split points for a given dimension takes  $N$  time once the data is sorted.

Without any optimisation, the complexity is  $\mathcal{O}(DN^2)$ . There are  $D$  dimensions and checking all the split points for a given dimension takes  $N^2$  time once the data is sorted.

## Answer 4

Following method trains a classification tree to predict outputs for each of the following possible depths: 1,3,6,9,12,14. Then using 5-fold cross validation out of sample classification error is determined.

These methods use `DecisionTreeClassifier` and `KFold` modules from `scikit-learn`.

```
1 import numpy as np
2 from sklearn.model_selection import KFold
3 from sklearn.tree import DecisionTreeClassifier
4
5 def train_classifier(X_trn, y_trn, depths):
6     errors = np.zeros((6, 2))
7     n_splits = 5
8     kf = KFold(n_splits=n_splits, random_state=None, shuffle=True)
9     for ind, depth in enumerate(depths):
10         clf = DecisionTreeClassifier(random_state=None, max_depth=
11             depth)
12         error = 0
13         for train_index, test_index in kf.split(X_trn):
14             # Split train-test
15             X_train, X_test = X_trn[train_index], X_trn[test_index]
16             y_train, y_test = y_trn[train_index], y_trn[test_index]
17
18             # Train the model
19             model = clf.fit(X_train, y_train)
20             predictions = model.predict(X_test)
21
22             error += 1 - sum([1 for actual, predicted in zip(y_test,
23             predictions) if actual==predicted])/len(X_test)
24             errors[ind] = depth, error/n_splits
25
26
27 data = np.load("data.npz")
28 X_trn = data["X_trn"]
29 y_trn = data["y_trn"]
30 X_tst = data["X_tst"]
31 train_classifier(X_trn, y_trn, {1, 3, 6, 9, 12, 14})
```

The below table lists the out of sample classification error for different depths.

<b>Depth</b>	<b>Mean classification error</b>
1	0.6365
3	0.54316667
6	0.498
9	0.50266667
12	0.51783333
14	0.50866667

Table 1: Mean out of sample classification error for Classification Tree

## Answer 5

1. We chose depth of tree to be **6**.
2. Our estimated Generalized error using 5-fold validation for this depth was **0.498**
3. On the public part of the leader-board the accuracy observed was **0.49722**

## Answer 6

The time complexity of KNN classifier that uses a brute force approach to pick K nearest points from N points is :

$$\mathcal{O}(ND + NK)$$

In a KNN classifier, the steps involved are as below :

1. Calculate distance of N points from the given point -  $\mathcal{O}(ND)$
2. Pick K smallest distances from N distances -  $\mathcal{O}(NK)$
3. Find label with max frequency in K outputs -  $\mathcal{O}(K)$

The over all time complexity -  $\mathcal{O}(ND + NK)$

## Answer 7

Following methods do nearest-neighbor prediction for each of the following possible values of K: 1, 3, 5, 7, 9, 11. Then using 5-fold cross validation out of sample classification error is determined.

These methods use `KNeighborsClassifier` and `KFold` modules from `scikit-learn`.

```
1 from preface import *
2
3 # KNN Classification
4 def KNN_classification(numNeighbors, X_trn, y_trn, X_val):
5     neigh = KNeighborsClassifier(n_neighbors=numNeighbors, n_jobs=-1)
6     neigh.fit(X_trn, y_trn)
7     y_pred = neigh.predict(X_val)
8     return y_pred
9
10 # Evaluate what capacity of KNN would be optimum
11 def KNNModelSelection() :
12     data = np.load("data.npz")
13     X_trn = data["X_trn"]
14     y_trn = data["y_trn"]
15     X_tst = data["X_tst"]
16
17     kf = KFold(n_splits=5,shuffle=True)
18     kf.get_n_splits(X_trn)
19
20     kFoldForTraining = []
21     kFoldForTesting = []
22
23     # Split and store the splits so that the same splits can be used
24     # for all values of K.
25     for train_index, test_index in kf.split(X_trn):
26         kFoldForTraining.append(train_index)
27         kFoldForTesting.append(test_index)
28
29     K = [1, 3, 5, 7, 9, 11]
30     classificationError = []
31     for i in K:
32         # For each of the splits, run KNN classification
33         classErrForThisValueOfK=0
34         for splitNum in range(len(kFoldForTraining)) :
35             X_train, X_test = X_trn[kFoldForTraining[splitNum]], X_trn[kFoldForTesting[splitNum]]
36             y_train, y_test = y_trn[kFoldForTraining[splitNum]], y_trn[kFoldForTesting[splitNum]]
37             y_pred = KNN_classification(i, X_train, y_train, X_test)
38             classErrForThisValueOfK = classErrForThisValueOfK + (1 - accuracy_score(y_test, y_pred))
39
40             # Store the mean classification error for this value of K in
41             # the numpy array
42             classificationError.append([i, classErrForThisValueOfK/5.0])
```

```

41     prettyPrintTable(classificationError, ["Num neighbors","Errors"])
42
43 if __name__ == '__main__':
44     KNNModelSelection()

```

The below table lists the out of sample classification error for different values of K.

K	Mean classification error
1	0.456667
3	0.453167
5	0.449000
7	0.439000
9	0.432667
11	0.432167

Table 2: Mean out of sample classification error for KNN classifier

## Answer 8

1. We chose K to be **11**.
2. Generalized error using 5-fold validation for this value of K was **0.432167**.
3. On the public part of the leader-board the accuracy observed was **0.61666**.

## Answer 9

The methods defined below perform logistic regression and hinge classification respectively for the given datasets with  $\lambda \in \{10^{-4}, 10^{-2}, 1, 10, 100\}$ , validate the results according to 5-fold cross validation , and print the required metrics (0-1 misclassification error, hinge loss, and logistic loss.)

These methods use LogisticRegression, LinearSVC and KFold modules from scikit-learn.

```
1 # Custom softmax function to obtain probabilities from confidence
2 # scores
3 def softmax(x):
4     return np.exp(x)/sum(np.exp(x))
5
6 # Function to return required mtrics (0-1 misclassification error,
7 # Logistic loss, Hinge Loss)
8 def getMetrics(l, kf, clf):
9     misc_error, ll_aggr, hl_aggr = 0, 0, 0
10    for trn_index, tst_index in kf.split(X_trn):
11        # Split training data accordingly to training and validation
12        # data sets
13        cv_x_trn, cv_x_tst = X_trn[trn_index], X_trn[tst_index]
14        cv_y_trn, cv_y_tst = y_trn[trn_index], y_trn[tst_index]
15
16        # Fit model to training set
17        model = clf.fit(cv_x_trn, cv_y_trn)
18
19        # Decision function to generate confidence values
20        decisions = model.decision_function(cv_x_tst)
21        cv_pred_tst = np.argmax(decisions, axis=1)
22
23        # Probability values for decisions
24        probs = softmax(decisions)
25
26        # Calculate miscalculation error for this particular fold
27        error_1 = 0
28        for x,y in zip(cv_y_tst, cv_pred_tst):
29            if x!=y:
30                error_1+=1
31        misc_error += error_1/len(cv_pred_tst)
32
33        # Calculate Log Loss for this particular fold
34        ll = log_loss(cv_y_tst, probs)
35        ll_aggr += ll
36
37        # Calculate Hinge Loss for this particular fold
38        hl = hinge_loss(cv_y_tst, decisions)
39        hl_aggr += hl
```

```

40     # Calculate average metrics across splits
41     avg_misc_error = misc_error/kf.n_splits
42     avg_ll = ll_aggr/kf.n_splits
43     avg_hl = hl_aggr/kf.n_splits
44
45     return avg_misc_error, avg_ll, avg_hl
46
47
48 # Validate Logistic Regression Model
49 def logistic_regressor():
50     l_vals=[0.0001, 0.01, 1, 10, 100]
51
52     # Define splits for K-Fold Cross-Validation
53     n_splits=5
54     # KFold Cross Validator
55     kf = KFold(n_splits=n_splits, shuffle=True)
56
56     misc_err = []
57     log_loss = []
58     hinge_loss = []
59
60     for l in l_vals:
61         log_reg = LogisticRegression(penalty='l2', C=1/(2*l), solver='sag',
62                                     max_iter=100, multi_class='multinomial',
63                                     n_jobs = -1)
64         me, ll, hl = getMetrics(l, kf, log_reg)
65         misc_err.append(me)
66         log_loss.append(ll)
67         hinge_loss.append(hl)
68
69     print("Lambda\t\tError Value")
70     for l in range(len(l_vals)):
71         print(f"{l_vals[l]}\t\t{misc_err[l]}")
72     print("\n")
73     print("Lambda\t\tLog Loss")
74     for l in range(len(l_vals)):
75         print(f"{l_vals[l]}\t\t{log_loss[l]}")
76     print("\n")
77     print("Lambda\t\tHinge Loss")
78     for l in range(len(l_vals)):
79         print(f"{l_vals[l]}\t\t{hinge_loss[l]}")
80
81
82 # Validate Hinge Classification Model
83 def hinge_classifier():
84     l_vals=[0.0001, 0.01, 1, 10, 100]
85
86     # Define splits for K-Fold Cross-Validation
87     n_splits=5
88     # KFold Cross Validator
89     kf = KFold(n_splits=n_splits, shuffle=True)
90

```

```

91     misc_err = []
92     log_loss = []
93     hinge_loss = []
94
95     for l in l_vals:
96         hinge_clf = LinearSVC(loss= 'hinge', C = 1/(2*l), max_iter
=1000)
97         me, ll, hl = getMetrics(l, kf, hinge_clf)
98         misc_err.append(me)
99         log_loss.append(ll)
100        hinge_loss.append(hl)
101
102    print("Lambda\t\tError Value")
103    for l in range(len(l_vals)):
104        print(f"{l_vals[l]}\t\t{misc_err[l]}")
105    print("\n")
106    print("Lambda\t\tLog Loss")
107    for l in range(len(l_vals)):
108        print(f"{l_vals[l]}\t\t{log_loss[l]}")
109    print("\n")
110    print("Lambda\t\tHinge Loss")
111    for l in range(len(l_vals)):
112        print(f"{l_vals[l]}\t\t{hinge_loss[l]}")
113
114
115 if __name__ == "__main__":
116     logistic_regressor()
117     hinge_classifier()

```

The below tables are the metrics observed for **logistic regression** model -

Lambda	Error Value
0.0001	0.4123333333333333
0.01	0.4121666666666667
1	0.41
10	0.3956666666666667
100	0.3866666666666666

Table 3: 0-1 Misclassification Error for Logistic Regression

Lambda	Log Loss
0.0001	1.9744296274108781
0.01	1.9387817406836756
1	1.8253285331018811
10	1.483022446742442
100	1.1412247852627455

Table 4: Log Loss for Logistic Regression

Lambda	Hinge Loss
0.0001	1.1703172713760126
0.01	1.1645862979908836
1	1.1242905902918365
10	0.9723185930244626
100	0.8656582863025036

Table 5: Hinge Loss for Logistic Regression

The below tables are the metrics observed for **hinge classification** model -

Lambda	Error Value
0.0001	0.4894999999999994
0.01	0.4941666666666664
1	0.4491666666666667
10	0.4073333333333333
100	0.3989999999999997

Table 6: 0-1 Misclassification Error for Hinge Classification

Lambda	Log Loss
0.0001	1.658265210815916
0.01	1.6599533809685176
1	1.5731331893304552
10	1.1616912901671272
100	1.0157286277097335

Table 7: Log Loss for Hinge Classification

Lambda	Hinge Loss
0.0001	1.4622453356275216
0.01	1.4936376941641398
1	1.3380753667748453
10	0.990812192361957
100	0.871463119505194

Table 8: Hinge Loss for Hinge Classification

## Answer 10

1. We chose to use logistic loss model, with  $\lambda$  set to be **100**.
2. Generalized error using 5-fold validation for this value of  $\lambda$  was **0.3867**. Log loss for this model was **1.1412**, and hinge loss for this model was **0.8657**.
3. On the public part of the leader-board, the accuracy observed was **0.62500**.

## Answer 11

```
1 def prediction_loss(x,y,W,V,b,c):
2     # Compute f(x) for all x.
3     Wx = np.matmul(W,x)
4     bPlusWx = np.add(b, Wx)
5
6     # fx = c + V.tanh(b + W.x)
7     fx = np.add(c, np.matmul(V, np.tanh(bPlusWx)))
8
9     expFx = 0
10    for fi in fx :
11        expFx = expFx + np.exp(fi)
12
13    # Loss = -f_y(x) + log(sum(exp(f_i(x))))
14    L = 0 - fx[y] + np.log(expFx)
15
16    return L
```

## Answer 12

```
1 def prediction_grad(x,y,W,V,b,c):
2     # Compute f(x) for all x. f(x) = c + V.tanh(b + W.x)
3     Wx = np.matmul(W,x)
4     bPlusWx = np.add(b, Wx)
5     fx = np.add(c, np.matmul(V, np.tanh(bPlusWx)))
6
7     # Compute dLdf = - e_hat_y + g(f(x)). g(f(x)) = exp(f_y) /
8     # sumOfAll(exp(f))
9     # e_hat_y is a unit vector with value 1 for y and 0 for rest.
10    # Compute negative unit vector
11    numLabels = c.shape[0]
12    e = np.zeros((numLabels, numLabels))
13    for i in range(numLabels) :
14        e[i][i] = -1
15
16    # Compute g(f(x))
17    sigma_exp_fv = 0
18    for i in range(numLabels) :
19        sigma_exp_fv = sigma_exp_fv + np.exp(fx[i])
20    gfx = np.zeros(numLabels)
21    for i in range(numLabels) :
22        gfx[i] = np.exp(fx[i]) / sigma_exp_fv
23
24    dLdf = e[y] + gfx
25
26    dLdc = dLdf
27
28    h = np.tanh(bPlusWx)
29    dLdV = np.outer(dLdf,h)
30
31    derivativeOfTanh = 1 - np.tanh(bPlusWx)**2
32    dLdb = derivativeOfTanh * np.matmul(np.transpose(V), dLdf)
33    dLdW = np.outer(dLdb,x)
34
35    return dLdW, dLdV, dLdb, dLdc
```

## Answer 13

$$1. \frac{dL}{dW} = \begin{pmatrix} -0.18070664 & -0.36141328 \\ -0.18070664 & -0.36141328 \\ 0.0 & 0.0 \end{pmatrix}$$

$$2. \frac{dL}{dV} = \begin{pmatrix} -0.45257413 & 0.45257413 & 0.48201379 \\ 0.45257413 & -0.45257413 & -0.48201379 \end{pmatrix}$$

$$3. \frac{dL}{db} = [ -0.18070664 \quad -0.18070664 \quad 0.0 ]$$

$$4. \frac{dL}{dc} = [ 0.5 \quad -0.5 ]$$

## Answer 14

```
1 from autograd import grad
2 from autograd import numpy as np
3
4 def prediction_loss(x,y,W,V,b,c):
5     Wx = np.matmul(W, x)
6     b_plus_Wx = np.add(b, Wx)
7     sigma = np.tanh(b_plus_Wx)
8     f = np.add(c, np.matmul(V, sigma))
9     softmax = 0
10    for item in f:
11        softmax += np.exp(item)
12
13    softmax = np.log(softmax)
14    return -f[y] + softmax
15
16 def prediction_grad_autograd(x,y,W,V,b,c):
17     dLdW = grad(prediction_loss, 2)(x,y,W,V,b,c)
18     dLdV = grad(prediction_loss, 3)(x,y,W,V,b,c)
19     dLdb = grad(prediction_loss, 4)(x,y,W,V,b,c)
20     dLdc = grad(prediction_loss, 5)(x,y,W,V,b,c)
21     return dLdW, dLdV, dLdb, dLdc
```

## Answer 15

```
1 from autograd import numpy as np
2
3 def prediction_loss_full(X,Y,W,V,b,c,l):
4     WX = np.matmul(W, np.transpose(X))
5     b = np.array([b, ] * X.shape[0]).transpose()
6     c = np.array([c, ] * X.shape[0]).transpose()
7     b_plus_WX = np.add(b, WX)
8     sigma = np.tanh(b_plus_WX)
9     f = np.add(c, np.matmul(V, sigma))
10
11    L = 0
12    for i in range(Y.shape[0]) :
13        softmax = np.log(np.sum(np.exp(f[:,i])))
14        y = Y[i]
15        L += -f[y][i] + softmax
16
17    L += l*(np.sum(np.square(V)) + np.sum(np.square(W)))
18
19    return L
```

## Answer 16

```
1 from autograd import grad
2
3 def prediction_grad_full(X,Y,W,V,b,c,l):
4     dLdW = grad(prediction_loss_full, 2)(X, Y, W, V, b, c, l)
5     dLdV = grad(prediction_loss_full, 3)(X, Y, W, V, b, c, l)
6     dLdb = grad(prediction_loss_full, 4)(X, Y, W, V, b, c, l)
7     dLdc = grad(prediction_loss_full, 5)(X, Y, W, V, b, c, l)
8
9     return dLdW, dLdV, dLdb, dLdc
```

## Answer 17

1. Below is the table describing the total training time (in ms) for all iterations.

M	Total time taken in Milliseconds
5	23919335
40	24029664
70	24103272

Table 9: Total time taken for 1000 iterations of gradient descent

2. Plot of regularized loss as a function of number of iterations -

